

Question 1: Skipped

A relational database is appropriate for which of the below scenarios?

- ☐ For those scenarios where there is a high volume of changes to relationships between entities
- ☐ For those scenarios where there is a high volume of geographic distributed writes
- ☒ For those scenarios where there is a high volume of transactional writes (Correct)
- ☐ For those scenarios where there is a high volume of writes that have varying data structures

Explanation

A relational database can be used for OLTP purposes. This is can be used for transactional systems. Hence if you have a high volume of transactional writes, you can make use of relational databases.

Option A is incorrect because if you gave a lot of relationship changes , then it would be an additional headache for the data in the table to adhere to those relationship changes.

Option B is incorrect because it becomes difficult to maintain the integrity of data if writes are happening from many locations.

Option D is incorrect because if there are varying data structures, it becomes difficult for the data to adhere to those varying data structures.

Question 2: Skipped

Which of the following is required for an extract, load and transform (ELT) process?

- ☐ A data pipeline that includes a transformation engine
- ☐ A separate transformation engine
- ☒ A target data store powerful enough to transform the data (Correct)
- ☐ Data that is fully processed before being loaded to the target data store

Explanation

In the ELT process, the transformation happens on the target system. Hence the target system must be powerful enough to transform the data.

Question 3: Skipped

A company wants to load data from a customer relationship management system onto a data warehouse by using an extract , load and transform process.

Which of the following is the stage in the ELT process you would perform the process of Extract?

☐ An in-memory data integration tool

☒ The CRM system

(Correct)

☐ The data warehouse

Explanation

Here since the data needs to be extracted from the CRM system, this will be the stage in the process of Extraction.

Question 4: Skipped

A company wants to load data from a customer relationship management system onto a data warehouse by using an extract , load and transform process.

Which of the following is the stage in the ELT process you would perform the process of Load?

☐ An in-memory data integration tool

☐ The CRM system

☒ The data warehouse

(Correct)

Explanation

Here since the data needs to be loaded into the data warehouse, this will be the stage for the Load process.

Question 5: Skipped

A company wants to load data from a customer relationship management system onto a data warehouse by using an extract , load and transform process.

Which of the following is the stage in the ELT process you would perform the process of Transform?

☒ An in-memory data integration tool (Correct)

☐ The CRM system

☐ The data warehouse

Explanation

Since the transformation is going to occur in the destination , that is the data warehouse, there needs to be an in-memory data integration tool to perform the transformation.



Question 6: Skipped

A company is planning on implementing the following workloads.

- 1) Storage of data for a product catalog that will be loaded every 12 hours to a data warehouse
- 2) Storage of data for online purchases will be loaded to a data warehouse as the purchases occur
- 3) Updates to the inventory data that will be loaded to a data warehouse every 1000 transactions

Which of the following is the workload type for the workload?

“Storage of data for a product catalog that will be loaded every 12 hours to a data warehouse”

☒ Batch

(Correct)

☐ Micro-Batch

☐ Streaming

Explanation

This is a batch process that is going to occur after a duration of time.

Question 7: Skipped

A company is planning on implementing the following workloads.

- 1) Storage of data for a product catalog that will be loaded every 12 hours to a data warehouse
- 2) Storage of data for online purchases will be loaded to a data warehouse as the purchases occur
- 3) Updates to the inventory data that will be loaded to a data warehouse every 1000 transactions

Which of the following is the workload type for the workload?

“Storage of data for online purchases will be loaded to a data warehouse as the purchases occur”

☐ Batch

☐ Micro-Batch

☒ Streaming

(Correct)

Explanation

Here since the data needs to be loaded as soon as the purchase occur, this is an example of the streaming workload type.

Question 8: Skipped

A company is planning on implementing the following workloads.

- 1) Storage of data for a product catalog that will be loaded every 12 hours to a data warehouse
- 2) Storage of data for online purchases will be loaded to a data warehouse as the purchases occur
- 3) Updates to the inventory data that will be loaded to a data warehouse every 1000 transactions

Which of the following is the workload type for the workload?

“Updates to the inventory data that will be loaded to a data warehouse every 1000 transactions”

☐ Batch

☒ Micro-Batch

(Correct)

☐ Streaming

Explanation

Here since the data need to be loaded into the data warehouse after a certain threshold, this is an example of the Micro-Batch workload type.

Question 9: Skipped

You need to match the right type of analytics that can be used for the following business scenario?

“Why did the sales increase last month”

☐ Cognitive

☐ Descriptive

☒ Diagnostics

(Correct)

☐ Predictive

☐ Prescriptive

Explanation

This is an example of Diagnostics analytics. This is used to understand why something happened.

Question 10: Skipped

You need to match the right type of analytics that can be used for the following business scenario?

“Which people are mentioned in a company’s business documents”

☐ Cognitive

☒ Descriptive

(Correct)

☐ Diagnostics

☐ Predictive

☐ Prescriptive

Explanation

This is an example of Descriptive analytics. Here you are describing information about the people mentioned in the business documents.

Question 11: Skipped

Which of the following is the right option when it comes to relational data?

☐ Relational data is stored in a file system as unstructured data

☐ Relational data is stored in a hierarchical folder structure

☒ Relational data is stored in a tabular form of rows and columns (Correct)

☐ Relational data is stored in a comma-separated value file

Explanation

Relational data is normally stored in the form of tables. The tables consist of rows and columns of data.

Question 12: Skipped

Which of the following are two uses of data visualization? Choose 2 answers from the options given

☐ Represent trends and patterns over time. (Correct)

☐ Communicate the significance of data. (Correct)

☐ Implement machine learning to predict future values.

☐ Consistently implement business logic across reports

Explanation

With data visualization, you can see patterns and trends in your data over time. You can also communicate the significance of data to your stakeholders.

Question 13: Skipped

A company is planning on implementing the following workloads.

- 1) Analyzing web traffic from the past year
- 2) Classifying images that were uploaded last month
- 3) Tracking how many people are currently using a website

Which of the following is the workload type for the workload?

“Analyzing web traffic from the past year”

☒ Batch

(Correct)

☐ Streaming

Explanation

Here since you are looking at taking data from a period of time, this qualifies as batch processing

Question 14: Skipped

A company is planning on implementing the following workloads.

- 1) Analyzing web traffic from the past year
- 2) Classifying images that were uploaded last month
- 3) Tracking how many people are currently using a website

Which of the following is the workload type for the workload?

“Classifying images that were uploaded last month”

☒ Batch

(Correct)

☐ Streaming

Explanation

Here since you are looking at taking data from a period of time, this qualifies as batch processing

Question 15: Skipped

A company is planning on implementing the following workloads.

- 1) Analyzing web traffic from the past year
- 2) Classifying images that were uploaded last month
- 3) Tracking how many people are currently using a website

Which of the following is the workload type for the workload?

“Tracking how many people are currently using a website”

☐ Batch


☒ Streaming

(Correct)

Explanation

Since here the user data needs to be tracked in real-time, this classifies as Streaming data.

You need to complete the architecture



Which of the following would go into Area 1?

☐ Azure Analysis services

☒ Azure Data Factory (Correct)

☐ Azure Cosmos DB

☐ Azure Synapse Analytics


☐ Azure Table storage

Explanation

Azure Data Factory can be used as the service to Extract , transform and load the data into the destination.

Question 17: Skipped

You need to complete the architecture



Which of the following would go into Area 2?

☐ Azure Analysis services

☐ Azure Data Factory

☐ Azure Cosmos DB

☒ Azure Synapse Analytics (Correct)

☐ Azure Table storage

Explanation

You can use Azure Synapse Analytics. Here you can host a SQL data warehouse.

Question 18: Skipped

You have an application that performs reads and writes data to an Azure SQL Database. Which of the following is the type of processing that can be performed on the database?

☐ Stream processing

☐ Batch processing

☐ Online Analytical Processing (OLAP)

☒ Online Transaction Processing (OLTP)

(Correct)

Explanation

Online Transaction Processing (OLTP) is performed on a relational database.

Question 19: Skipped

Which of the following are two aspects of real-time data processing?

☐ Data is processed as it is created.

(Correct)

☐ Low latency is expected

(Correct)

☐ High latency is expected

☐ Data is processed periodically

Explanation

Here data needs to be processed immediately. And here low latency is expected. You should be able to retrieve the data immediately so that it can be processed.

Question 20: Skipped

Which of the following are two Azure services that can be used to provision Spark clusters?

☒ Azure Databricks

(Correct)

☐ Azure Log Analytics

☐ Azure Time Series Insights

☒ Azure HDInsight

(Correct)

Explanation

You can create Spark clusters using Azure Databricks and Azure HDInsight.

Question 21: Skipped

You have to state whether the following is TRUE or FALSE.

“A pipeline is a representation of a data structure within Azure Data Factory”

☐ TRUE

☒ FALSE

(Correct)

Explanation

The data set in Azure Data Factory is used to represent the data structure.

Question 22: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Data Factory can execute other pipelines”

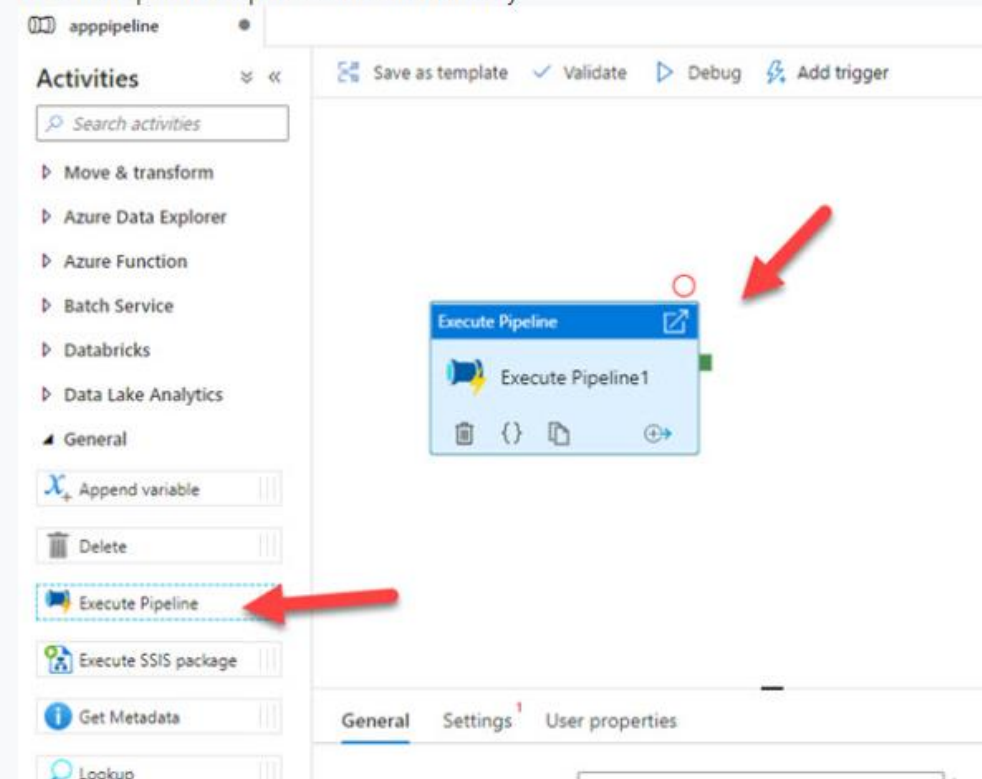
☐ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can trigger other pipelines from Azure Data Factory. You can use the Execute Pipeline step in Azure Data Factory.



Question 23: Skipped

You have to state whether the following is TRUE or FALSE.

“A processing step within an Azure Data Factory pipeline is an activity”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, activities in Azure Data Factory can be used as a processing step to process the data.

Question 24: Skipped

You have to work with the Azure Data Factory service. You need to map the Azure Data Factory components to the right descriptions.

Which of the following component should be mapped to the following description?

“ A representation of data structures within data stores”

☒ Dataset

(Correct)

☐ Linked service

☐ Mapping data flow

☐ Pipeline

Explanation

The Dataset is a representation of the data structured within data stores.

Question 25: Skipped

Which of the following component should be mapped to the following description?

“The information used to connect to external data sources”

☐ Dataset

☒ Linked service

(Correct)

☐ Mapping data flow

☐ Pipeline



Explanation

The Linked Service can be used to connect to external data sources.

Question 26: Skipped

Which of the following component should be mapped to the following description?

“ A logical grouping of activities that performs a unit of work and can be scheduled”

☐ Dataset

☐ Linked service

☐ Mapping data flow

☒ Pipeline

(Correct)

Explanation

Here you can define a pipeline of activities

Question 27: Skipped

You need to select the appropriate an Azure service that would be used for a particular requirement. Which of the following would you use for the following requirement?

“Output data to Parquet format”

☒ Azure Data Factory (Correct)

☐ Azure Data Lake Storage

☐ Azure SQL Database

☐ Azure Synapse Analytics

Explanation

You can use Azure Data Factory to output data in a particular format.

Question 28: Skipped

Which of the following would you map to the Graph data structure?

☐ Cassandra API

☒ Gremlin API

(Correct)

☐ MongoDB API

☐ Table API

Explanation

The Gremlin API supports the usage of graph database.

Question 29: Skipped

You need to map the appropriate Azure Cosmos DB API's to the appropriate data structures

Which of the following would you map to the JSON data structure?

☐ Cassandra API

☐ Gremlin API

☒ MongoDB API

(Correct)

☐ Table API

Explanation

The MongoDB API supports the usage of documents that can be used to store JSON data.

Question 30: Skipped

You need to map the appropriate Azure Cosmos DB API's to the appropriate data structures

Which of the following would you map to the Key/Value data structure?

☐ Cassandra API

☐ Gremlin API

☐ MongoDB API

☒ Table API

(Correct)

Explanation

The Table API supports the usage of key/value pairs.

Question 31: Skipped

You currently manage an application that stores data within a shared folder on a Windows server. You have to move the shared folder to Azure Storage. Which of the following service within an Azure storage account could you use for this purpose?

☐ Queue

☐ Blob

☒ File

(Correct)

☐ Table

Explanation

You can use the File service to create file shares. These file shares can be mounted on Windows Servers.

Question 32: Skipped

You have to state whether the following is TRUE or FALSE.

“When data is ingested from Azure Data Lake Storage across Azure regions, then you will incur costs for bandwidth”

☒ TRUE

(Correct)

☐ FALSE

Explanation

For any data transfer costs across regions, there is a cost associated.

Question 33: Skipped

You have to state whether the following is TRUE or FALSE.

“You can use the blob, table and file storage service in the same Azure Storage account”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can use these services at the same time in an Azure Storage account.

Question 34: Skipped

You have to state whether the following is TRUE or FALSE.

“You can implement Azure Data Lake Storage by creating an Azure Storage account”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, since Azure Data Lake Storage is built on Azure Storage Accounts

Question 35: Skipped

Which of the following is an Azure Data Factory component that provides the compute environment for activities?

☐ A linked service

☒ An integration runtime

(Correct)

☐ A control flow

☐ A pipeline

Explanation

This is provided by the integration runtime.

Question 36: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Databricks is an Apache Spark-based collaborative analytics platform”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Azure Databricks is an Apache Spark based platform.

Question 37: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Data Factory orchestrates data Integration workflows”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Azure Data Factory is used to create data integration workflows.

Question 38: Skipped

You need to select the appropriate an Azure service that would be used for a particular requirement. Which of the following would you use for the following requirement?

“Store data in Parquet format”

☐ Azure Data Factory

☒ Azure Data Lake Storage

(Correct)

☐ Azure SQL Database

☐ Azure Synapse Analytics

Explanation

You can store files in virtually any data format in Azure Data Lake Storage.

Question 39: Skipped

You need to select the appropriate an Azure service that would be used for a particular requirement. Which of the following would you use for the following requirement?

“Persist a tabular representation of data that is stored in Parquet format”

☐ Azure Data Factory

☐ Azure Data Lake Storage

☐ Azure SQL Database

☒ Azure Synapse Analytics

(Correct)

Explanation

The tabular representation of the data is normally stored in Azure Synapse Analytics.

Question 40: Skipped

Which of the following is the right statement when it comes to the massively parallel processing engine in Azure Synapse Analytics?

- ☒ The massively parallel processing engine of Azure Synapse Analytics distributes processing across compute nodes (Correct)
- ☐ The massively parallel processing engine of Azure Synapse Analytics distributes processing across control nodes
- ☐ The massively parallel processing engine of Azure Synapse Analytics redirects client connections across compute nodes
- ☐ The massively parallel processing engine of Azure Synapse Analytics redirects client connections across control nodes

Explanation

The engine actually sends all the processing to the compute nodes.

Question 41: Skipped

Which of the following is an Azure storage solution that provides native support for POSIX-compliant access control lists?

☐ Azure Queue storage

☒ Azure Data Lake Storage

(Correct)

☐ Azure table storage

☐ Azure Files

Explanation

This feature is provided by Azure Data Lake Storage

Question 42: Skipped

Your company currently has a transactional application that stores data in an Azure SQL Managed instance. In which of the following circumstances would you need to implement a read-only database replica?

- ☒ You need to generate reports without affecting the transactional workload. (Correct)
- ☐ You need to audit the transactional application.
- ☐ You need to implement high availability in the event of a regional outage
- ☐ You need to improve the recovery point objective (RPO).

Explanation

A read-only instance can be used to reduce the workload on your OLTP database. You can generate reports from the read-only copy of the database.

Question 43: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Synapse Analytics scales storage and compute independently”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, these are two separate aspects when it comes to Azure Synapse Analytics

Question 44: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Synapse Analytics can be paused to save on compute costs”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can pause the Synapse pool to save on costs.

Question 45: Skipped

You have to state whether the following is TRUE or FALSE.

“An Azure Synapse Analytics data warehouse has fixed storage capacity”

☐ TRUE

☒ FALSE

(Correct)

Explanation

For Gen2 data warehouse, you have unlimited storage for columnstore tables.

Question 46: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Databricks can consume data from an Azure SQL database”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can consume data from an Azure SQL database

Question 47: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Databricks can consume data from Azure Event Hubs”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can consume data from an Azure Event Hubs

Question 48: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure Databricks can consume data from Azure Cosmos DB”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can consume data from an Azure Cosmos DB

Question 49: Skipped

Which of the following is TRUE when it comes a key/value data store?

☐ A key value data store is optimized for enforcing constraints

☒ A key value data store is optimized for simple lookups (Correct)

☐ A key value data store is optimized for table joins 

☐ A key value data store is optimized for transactions

Explanation

A key value store is better for simple lookups based on the key.

Question 50: Skipped

Your company has an application that runs on Windows. It needs access to a mapped drive. Which of the following is an Azure service you can use for this requirement?

☐ Azure Cosmos DB

☐ Azure Table storage

☒ Azure Files

(Correct)

☐ Azure Blob storage

Explanation

Azure Files allows you to make use of mapped drives

Question 51: Skipped

Which of the following is TRUE when it comes a graph database in Azure Cosmos DB?

- ☐ You can query a graph database in Azure Cosmos DB as a JSON document by using a SQL-like language
- ☐ You can query a graph database in Azure Cosmos DB as a partitioned row store by using Cassandra Query language
- ☐ You can query a graph database in Azure Cosmos DB as a partitioned row store by using Language-Integrated Query
- ☒ You can query a graph database in Azure Cosmos DB as nodes and edges by using Gremlin language (Correct)

Explanation

The graph database is based on nodes and edges. And you can use the Gremlin API in Azure Cosmos DB to host a graph based database.

Question 52: Skipped

Your company needs to design a database that shows how traffic changes in one area of a network affect other components on the network. Which of the following is a data store type you would use for this requirement?

☐ Key/value

☒ Graph

(Correct)

☐ Document

☐ Columnar

Explanation

It would be ideal to define the network as nodes in a graph-based database. You can define the relationship between the networks as edges between the nodes.

Question 53: Skipped

Your company needs a non-relational data store that is optimized for storing and retrieving files, videos, audio streams and virtual disk images. The data store must be able to store data , metadata and a unique ID for each file.

Which of the following is a data store type you would choose for this requirement?

☐ Document

☐ Key/value

☒ Object

(Correct)

☐ Columnar

Explanation

For this you can use an Object data store like Azure Blob Storage. Here you can have metadata for each object and also a unique URL.

Question 54: Skipped

Your company wants to have an Azure storage account in place. They want to ensure that the data is replicated in another Azure region automatically. Which of the following can be used as the replication strategy for the storage account?

Choose 2 answers from the options given below

☐ Read-access geo-redundant storage (RA_GRS)

(Correct)

☐ Zone-redundant storage (ZRS)

☐ Geo-redundant storage (GRS)

(Correct)

☐ Locally-redundant storage (LRS)

Explanation

For this you can make use of Read-access geo-redundant storage (RA_GRS) and Geo-redundant storage (GRS). This ensures that data is replicated to another region.

Question 55: Skipped

You have to state whether the following is TRUE or FALSE.

“The Azure Cosmos DB API is configured separately for each database in an Azure Cosmos DB account”

☐ TRUE

☒ FALSE

(Correct)

Explanation

The Cosmos DB API is configured at the account level and not at the database level.

Question 56: Skipped

You have to state whether the following is TRUE or FALSE.

“Partition keys are used in Azure Cosmos DB to optimize queries”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, the data in an Azure Cosmos DB database is partitioned via the partition keys.

Question 57: Skipped

You have to state whether the following is TRUE or FALSE.

“Items contained in the same Azure Cosmos DB logical partition can have different partition keys”

☐ TRUE

☒ FALSE

(Correct)

Explanation

All items in the same logical partition have the same partition key.

Question 58: Skipped

Which of the following is the right use case for an Azure Resource Manager template?

- ☒ To automate the creation of an interdependent group of Azure resources in a repeatable way (Correct)
- ☐ To automate the creation of an interdependent resources that can be deployed from the Azure portal
- ☐ To provision Azure subscriptions and apply policies for multi-tenant deployments
- ☐ To set complex permissions automatically on existing objects in an Azure subscription

Explanation

Azure Resource Manager templates can be used to deploy resources to Azure. You can reuse the same templates to deploy the same set of resources in a repeatable way.

Question 59: Skipped

Which of the following is a feature in Azure Cosmos DB that provides redundancy within an Azure region?

☐ Multi-master replication

☒ Availability Zones

(Correct)

☐ Automatic failover

☐ The strong consistency level

Explanation

When Availability Zones, data is copied onto multiple Availability Zones. Even if one zone goes down, you still have the other zones available.

Option A is incorrect because this is an option when you enable writes onto multiple regions for Azure Cosmos DB.

Option C is incorrect because automatic failover happens when data is replicated across multiple regions

Option D is incorrect because consistency is an option for maintain the consistency of data when you have multiple regions enabled for your Azure Cosmos DB account.

Question 60: Skipped

At which of the following level in Azure Cosmos DB can you configure multiple write regions and read regions?

☐ database

☐ partition

☐ collection

☒ account

(Correct)

Explanation

This is done at the account level.

Question 61: Skipped

Your company is planning on designing an application that will write a high volume of JSON data that will have an application-defined schema. Which of the following can you specify as the data store type for this requirement?

☐ key/value

☒ document

(Correct)

☐ graph

☐ columnar

Explanation

For this you can use a document-based data store. Here the JSON data can be stored as documents.

Question 62: Skipped

At which of the following levels can you set the throughput for an Azure Cosmos DB account? Choose 2 answers from the options given below

☒ container

(Correct)

☐ item

☒ database

(Correct)

☐ partition

Explanation

You can define the throughput at a database or container level.

Question 63: Skipped

You have the following table in place

ProductName	Quantity
ProductA	100
ProductB	200
ProductC	300

Which statement should you use in a SQL query to change the quantity of ProductB to 400?

☐ INSERT

☐ MERGE

☒ UPDATE

(Correct)

☐ CREATE

Explanation

You would use the UPDATE statement for this requirement

Question 64: Skipped

Your company needs to implement a relational database in Azure. The solution must minimize ongoing maintenance. Which of the following can be implemented for this requirement?

☐ Azure HDInsight

☒ Azure SQL Database

(Correct)

☐ Azure Cosmos DB

☐ SQL Server on Azure virtual machines

Explanation

With the Azure SQL database, you don't need to maintain the underlying infrastructure.

Question 65: Skipped

You have to write a set of queries that will help administrators to troubleshoot an Azure SQL database. You have to be able to embed the documents and query results into a SQL notebook. Which of the following would you use for this requirement?

☐ Microsoft SQL Server Management Studio (SSMS)

☒ Azure Data Studio

(Correct)

☐ Azure CLI

☐ Azure PowerShell

Explanation

You can use SQL Notebooks in Azure Data Studio.

Question 66: Skipped

You have to map the right term to the term description.

Which of the following would you map to the following description?

“ A database object that holds data”

☐ Index

☐ View

☒ Table

(Correct)

Explanation

The table is used to hold data

Question 67: Skipped

You have to map the right term to the term description.

Which of the following would you map to the following description?

“ A database object whose content is defined by a query”

☐ Index

☒ View

(Correct)

☐ Table

Explanation

The View is defined by a query.



Question 68: Skipped

You have to map the right term to the term description.

Which of the following would you map to the following description?

“ A database object that helps improve the speed of data retrieval”

☒ Index

(Correct)

☐ View

☐ Table

Explanation

The index defined on the table can be used to improve the speed of data retrieval.

Question 69: Skipped

Which of the following is a command line tool that can be used to query Azure SQL databases?

☒ sqlcmd

(Correct)

☐ bcp

☐ azdata

☐ Azure CLI

Explanation

The sqlcmd tool can be used to query Azure SQL databases.

Question 70: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure SQL Database includes a fully managed backup service”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, there is a fully managed backup service that takes automated backups.

Question 71: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure SQL Database has built-in high availability”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, Azure SQL Database automatically has high availability,

Question 72: Skipped

You have to state whether the following is TRUE or FALSE.

“Azure SQL Database can use Azure Advanced Threat Protection”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, Azure SQL Database also can make use of Advanced Threat Protection.

Question 73: Skipped

You have to state whether the following is TRUE or FALSE.

“You can use Azure Data Studio to query a Microsoft SQL Server big data cluster”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can use Azure Data Studio to connect to a Microsoft SQL Server Big Data cluster

Question 74: Skipped

You have to state whether the following is TRUE or FALSE.

“You can use Microsoft SQL Server Management Studio to query an Azure Synapse Analytics data warehouse”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can use Microsoft SQL Server Management Studio to query a data warehouse hosted in Azure Synapse Analytics.

Question 75: Skipped

You have to state whether the following is TRUE or FALSE.

“You can use MySQL Workbench to query Azure Database for MariaDB databases”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes, you can use MySQL Workbench to query Azure Database for MariaDB.

Question 76: Skipped

You have to state whether the following is TRUE or FALSE.

“You must apply patches to Azure SQL databases regularly”

☐ TRUE

☒ FALSE

(Correct)

Explanation

The patches are applied automatically by the Azure SQL database service.

Question 77: Skipped

You have to state whether the following is TRUE or FALSE.

“You need a Microsoft 365 subscription to create an Azure SQL database”

☐ TRUE

☒ FALSE

(Correct)

Explanation

You need an Azure subscription to create an Azure SQL database.

Question 78: Skipped

Your company has to design a data store that would contain student data. A sample of the data is shown below

StudentNumber	StudentInformation
123344	FirstName: James Last:Smith
345554	FirstName: John Last: Dean Email: john@gmail.com
325456	Phone number: 122-900-1101

Which of the following would be the ideal data store type?

☐ Graph

☐ Key/value

☐ Object

☒ Columnar

(Correct)

Explanation

Here you can store the data as a column-based data store

One can refer to the below URL for the different types of data stores.

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

Question 79: Skipped

Which storage solution supports role-based access control (RBAC) at the file and folder level?

☐ Azure Disk Storage

☒ Azure Data Lake Storage

(Correct)

☐ Azure Blob storage

☐ Azure Queue storage

Explanation

This is possible with Azure Data Lake Storage

Question 80: Skipped

Your company need to store data in Azure Blob storage. The data needs to be stored for seven years. The retrieval time of the data is unimportant. The solution must minimize storage costs. Which of the following is the ideal storage tier to use for this requirement?

☒ Archive

(Correct)

☐ Hot

☐ Cool

Explanation

You can use the Archive tier for this requirement.

Question 81: Skipped

Which of the following is an example of an analytics bar chart that shows year-to-date sales by region?

☒ descriptive

(Correct)

☐ diagnostic

☐ predictive


☐ prescriptive

Explanation

This should be an example on descriptive analysis.

Question 82: Skipped

Your team currently has the following chart



What is the type of analysis is being depicted in the 4th quarter?

☐ Predictive

(Correct)

☐ Prescriptive

☐ Descriptive

☐ Diagnostic

Explanation

This is known as predictive analysis on what is going to happen based on past data.

Question 83: Skipped

You need to match the right type of analytics that can be used for the following business scenario?

“How do I allocate my budget to buy different inventory items”

☐ Cognitive

☐ Descriptive

☐ Diagnostics

☐ Predictive

☒ Prescriptive

(Correct)

Explanation

This should fall under prescriptive analysis. Here you want to know how to approach to fulfil a target requirement based on the data that is available.

Question 84: Skipped

Which of the following statement is right when it comes to Microsoft Power BI?

- ☐ A Microsoft Power BI dashboard enables users to create highly formatted ,fixed layout documents
- ☐ A Microsoft Power BI interactive report enables users to create highly formatted ,fixed layout documents
- ☒ A Microsoft Power BI paginated report enables users to create highly formatted ,fixed layout documents (Correct)

Explanation

With paginated reports, you can create highly formatted and fixed layout documents

Question 85: Skipped

You have to develop a solution that needs to provide data to senior executives. Below are the key requirements

- a) The solution must provide an interactive graphic interface
- b) It must depict various key performance indicators
- c) It must support data exploration by using drill down

Which of the following can be used within Microsoft BI for this requirement?

☒ A dashboard

(Correct)

☐ Microsoft Power Apps

☐ A dataflow

☐ A report

Explanation

For creating solutions for senior management, you would go ahead and make use of dashboards. Here you can add key performance indicators, provides a graphic interface and you can also perform drill down as well.

Question 86: Skipped

Which of the following needs to be used to build a Microsoft Power BI paginated report?

☒ Power BI Report Builder (Correct)

☐ Power BI Desktop

☐ Charticator

☐ The Power BI service

Explanation

To build a paginated report, you need to use the PowerBI Report builder.

Question 87: Skipped

You have to develop a reporting solution for your company. The solution needs to support paginated reports. The reports should be able to query a dimensional model in a data warehouse. Which of the following type of processing must the reporting solution use?

☐ A stream processing

☐ Online Transaction Processing (OLTP)

☐ Batch processing

☒ Online Analytical Processing (OLAP)

(Correct)

Explanation

Since we are connecting to a data warehouse, this needs to be an Online Analytical Processing system.

Question 88: Skipped

Which of the following are activities that are performed in the Microsoft Power BI service? Choose 2 answers from the options given below

☐ Report and dashboard creation

(Correct)

☐ Report sharing and distribution

(Correct)

☐ Data modeling

☐ Data acquisition and preparation

Explanation

Here you can create your reports and dashboards. You can also share reports with other users

Question 89: Skipped

You have to state whether the following is TRUE or FALSE.

“A Microsoft Power BI dashboard is associated with a single workspace”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes , the workspace is a logical boundary for your reports and data

Question 90: Skipped

You have to state whether the following is TRUE or FALSE.

“A Microsoft Power BI dashboard can only display visualizations from a single dataset”

☐ TRUE

☒ FALSE

(Correct)

Explanation

A dashboard can use multiple data sets for visualization.

Question 91: Skipped

You have to state whether the following is TRUE or FALSE.

“A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook”

☒ TRUE

(Correct)

☐ FALSE

Explanation

Yes , data can be visualized from Microsoft Excel